

Updates and Imports

```
!sudo apt update
!sudo apt autoremove
!sudo apt upgrade
!sudo apt install python3x

!python --version
```

```
import sys
print("version:", sys.version)
```

```
!pip install -U pandas
```

```
Looking in indexes: https://pypi.org/simple, https://us-
python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: pandas in
/usr/local/lib/python3.8/dist-packages (1.3.5)
Collecting pandas
  Downloading pandas-1.5.2-cp38-cp38-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (12.2 MB)
ent already satisfied: python-dateutil>=2.8.1 in
/usr/local/lib/python3.8/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: numpy>=1.20.3 in
/usr/local/lib/python3.8/dist-packages (from pandas) (1.21.6)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.8/dist-packages (from pandas) (2022.6)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.8/dist-packages (from python-dateutil>=2.8.1-
>pandas) (1.15.0)
Installing collected packages: pandas
  Attempting uninstall: pandas
    Found existing installation: pandas 1.3.5
    Uninstalling pandas-1.3.5:
      Successfully uninstalled pandas-1.3.5
Successfully installed pandas-1.5.2
```

```
{"pip_warning":{"packages":["pandas"]}}
```

```
from google.colab import drive
drive.mount('/content/drive/')
```

Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount("/content/drive/", force_remount=True).

```
import seaborn as sns
import cv2
import pandas as pd
import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.layers import Conv2D, MaxPooling2D,
    BatchNormalization
from tensorflow.keras.layers import Dense, Dropout, Flatten,
    Activation
from tensorflow.keras.metrics import categorical_accuracy,
    top_k_categorical_accuracy, categorical_crossentropy
from tensorflow.keras.models import Sequential
from tensorflow.keras.callbacks import EarlyStopping,
    ReduceLROnPlateau, ModelCheckpoint
from tensorflow.keras.optimizers import Adam
from tensorflow.keras import regularizers
from tensorflow.keras.applications import MobileNet, MobileNetV2
from tensorflow.keras.applications.mobilenet import preprocess_input
import matplotlib.pyplot as plt
from tensorflow.keras.utils import to_categorical
```

```
import pandas as pd
import numpy as np
from numpy import newaxis
from sklearn.model_selection import train_test_split
import pickle
from tqdm.notebook import tqdm
import matplotlib.pyplot as plt
import cv2
```

```
from numpy import newaxis
from sklearn.model_selection import train_test_split
import pickle
from tqdm.notebook import tqdm
import matplotlib.pyplot as plt
import cv2
```

```

train_path =
    "/content/drive/MyDrive/AMLminiproject/train100c5k_v2.pkl"
test_path =
    "/content/drive/MyDrive/AMLminiproject/test100c5k_nolabel.pkl"
# train_data = open(train_path, 'rb')
# test_data = open(test_path, 'rb')

# df_train = pd.read_pickle(train_path)
# df_test = pd.read_pickle(test_path)
df_train = pickle.load(open(train_path, 'rb'))
df_test = pickle.load(open(test_path, 'rb'))
# train_data = df['data'].values

```

Mobilenet

reference: <https://www.kaggle.com/code/sabarish2611/alexnet-vs-mobilenet-using-mnist-data#AlexNet-vs-MobileNet-:-Comparing-the-performance>

```

# sampling
sample = df_train.groupby('target', group_keys=False).apply(lambda x:
    x.sample(frac=0.3))
sample['data'] = sample['data']/255


X = np.array(list(sample['data'].values))
# X = np.array([img[:, :, newaxis] for img in
    np.array(sample['data'].values)])
y = np.array(sample['target'].values)
df_train = None
df_test = None
sample = None
X.shape

(150000, 28, 28)

plt.imshow(X[0])

<matplotlib.image.AxesImage at 0x7f8cd35c4d30>

```



```
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    random_state=42)
```

```
X = None
```

```
y = None
```

```
X_train.shape
```

```
(112500, 28, 28)
```

```
# X_re = np.zeros((X.shape[0],32,32,3))
# for i in tqdm(range(X.shape[0])):
#     temp = cv2.resize(X[i], dsize=(32, 32),
#                       interpolation=cv2.INTER_CUBIC)
#     # temp = temp[:, :, newaxis]
#     # print(temp.shape)
#     img = np.zeros((32,32,3))

#     # print(img.shape)
#     img[:, :, 0] = temp
#     img[:, :, 1] = temp
#     img[:, :, 2] = temp
#     X_re[i] = img
X_train_re = np.zeros((X_train.shape[0],32,32,3))
for i in tqdm(range(X_train.shape[0])):
    temp = cv2.resize(X_train[i], dsize=(32, 32),
                      interpolation=cv2.INTER_CUBIC)
    # temp = temp[:, :, newaxis]
    # print(temp.shape)
    img = np.zeros((32,32,3))

    # print(img.shape)
    img[:, :, 0] = temp
```

```

        img[:, :, 1] = temp
        img[:, :, 2] = temp
        X_train_re[i] = img
        X_train[i] = None
X_train = None
X_test_re = np.zeros((X_test.shape[0], 32, 32, 3))
for i in tqdm(range(X_test.shape[0])):
    temp = cv2.resize(X_test[i], dsize=(32, 32),
                      interpolation=cv2.INTER_CUBIC)
    # temp = temp[:, :, newaxis]
    # print(temp.shape)
    img = np.zeros((32, 32, 3))

    # print(img.shape)
    img[:, :, 0] = temp
    img[:, :, 1] = temp
    img[:, :, 2] = temp
    X_test_re[i] = img
    X_test[i] = None

X_test = None

{"version_major":2,"version_minor":0,"model_id":"5129e5f8573d4a298f261c9ecc7e3ea7"}

{"version_major":2,"version_minor":0,"model_id":"7a4f5d610cf3456fbbf2914b758803cb"}

X_train_re.shape

(112500, 32, 32, 3)

base_model = MobileNet(include_top=False,
                        weights='imagenet', input_shape = (32, 32, 3), classes=100)
base_model.trainable = True

for layer in base_model.layers[:50]:
    layer.trainable = False

MobileNet_model = Sequential()
MobileNet_model.add(base_model)
MobileNet_model.add(Flatten())
MobileNet_model.add(Dense(100, activation=('softmax')))
```

```
early_stopping = EarlyStopping(min_delta = 0.001, patience =
    20, restore_best_weights = True, verbose = 0)

# Compile
MobileNet_model.compile(optimizer = "adam" , loss =
    'sparse_categorical_crossentropy' , metrics = ['accuracy'])
```

```
# Train
Mobile = MobileNet_model.fit(X_train_re, y_train, batch_size = 400,
    epochs = 50, callbacks = [early_stopping], validation_data =
    (X_test_re, y_test))
```

```
MobileNet_model.summary()
```

WARNING:tensorflow: `input_shape` is undefined or non-square, or `rows` is not in [128, 160, 192, 224]. Weights for input shape (224, 224) will be loaded as the default.

Epoch 1/50

282/282 [=====] - 22s 34ms/step - loss: 2.3421 - accuracy: 0.4379 - val_loss: 2.2859 - val_accuracy: 0.4241

Epoch 2/50

282/282 [=====] - 7s 26ms/step - loss: 1.5940 - accuracy: 0.5846 - val_loss: 1.7769 - val_accuracy: 0.5469

Epoch 3/50

282/282 [=====] - 8s 28ms/step - loss: 1.3610 - accuracy: 0.6362 - val_loss: 1.7431 - val_accuracy: 0.5750

Epoch 4/50

282/282 [=====] - 7s 27ms/step - loss: 1.2019 - accuracy: 0.6715 - val_loss: 1.6661 - val_accuracy: 0.5821

Epoch 5/50

282/282 [=====] - 7s 26ms/step - loss: 1.0716 - accuracy: 0.7021 - val_loss: 1.7595 - val_accuracy: 0.5697

Epoch 6/50

282/282 [=====] - 7s 26ms/step - loss: 0.9593 - accuracy: 0.7280 - val_loss: 1.7388 - val_accuracy: 0.5848

Epoch 7/50

282/282 [=====] - 7s 26ms/step - loss: 0.8498 - accuracy: 0.7543 - val_loss: 1.8712 - val_accuracy: 0.5755

Epoch 8/50

282/282 [=====] - 7s 26ms/step - loss: 0.7601 - accuracy: 0.7752 - val_loss: 1.8884 - val_accuracy: 0.5788

Epoch 9/50

282/282 [=====] - 8s 29ms/step - loss: 0.6592
- accuracy: 0.8015 - val_loss: 2.0165 - val_accuracy: 0.5817
Epoch 10/50
282/282 [=====] - 9s 32ms/step - loss: 0.5815
- accuracy: 0.8218 - val_loss: 2.1094 - val_accuracy: 0.5749
Epoch 11/50
282/282 [=====] - 9s 32ms/step - loss: 0.5015
- accuracy: 0.8446 - val_loss: 2.2363 - val_accuracy: 0.5777
Epoch 12/50
282/282 [=====] - 8s 29ms/step - loss: 0.4281
- accuracy: 0.8653 - val_loss: 2.2120 - val_accuracy: 0.5758
Epoch 13/50
282/282 [=====] - 7s 26ms/step - loss: 0.3791
- accuracy: 0.8791 - val_loss: 2.3544 - val_accuracy: 0.5704
Epoch 14/50
282/282 [=====] - 7s 26ms/step - loss: 0.3340
- accuracy: 0.8926 - val_loss: 2.3482 - val_accuracy: 0.5773
Epoch 15/50
282/282 [=====] - 7s 26ms/step - loss: 0.2903
- accuracy: 0.9072 - val_loss: 2.4643 - val_accuracy: 0.5824
Epoch 16/50
282/282 [=====] - 8s 27ms/step - loss: 0.2675
- accuracy: 0.9130 - val_loss: 2.4808 - val_accuracy: 0.5766
Epoch 17/50
282/282 [=====] - 7s 26ms/step - loss: 0.2334
- accuracy: 0.9242 - val_loss: 2.5780 - val_accuracy: 0.5808
Epoch 18/50
282/282 [=====] - 8s 28ms/step - loss: 0.2231
- accuracy: 0.9278 - val_loss: 2.6697 - val_accuracy: 0.5767
Epoch 19/50
282/282 [=====] - 7s 26ms/step - loss: 0.2066
- accuracy: 0.9327 - val_loss: 2.6816 - val_accuracy: 0.5785
Epoch 20/50
282/282 [=====] - 7s 26ms/step - loss: 0.1954
- accuracy: 0.9373 - val_loss: 2.7885 - val_accuracy: 0.5827
Epoch 21/50
282/282 [=====] - 7s 26ms/step - loss: 0.1681
- accuracy: 0.9448 - val_loss: 2.7846 - val_accuracy: 0.5888
Epoch 22/50
282/282 [=====] - 7s 26ms/step - loss: 0.1701

```

- accuracy: 0.9436 - val_loss: 2.7298 - val_accuracy: 0.5878
Epoch 23/50
282/282 [=====] - 7s 26ms/step - loss: 0.1674
- accuracy: 0.9451 - val_loss: 2.8234 - val_accuracy: 0.5800
Epoch 24/50
282/282 [=====] - 8s 27ms/step - loss: 0.1504
- accuracy: 0.9499 - val_loss: 2.8720 - val_accuracy: 0.5853
Model: "sequential"

```

Layer (type)	Output Shape	Param #
=====		
mobilenet_1.00_224 (Functional)	(None, 1, 1, 1024)	3228864
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 100)	102500

```

=====
Total params: 3,331,364
Trainable params: 2,766,948
Non-trainable params: 564,416

```

```

base_model = MobileNet(include_top=False,
                        weights='imagenet', input_shape = (32,32,3), classes=100)
base_model.trainable = True

```

```

for layer in base_model.layers[:50]:
    layer.trainable = False

```

```

MobileNet_model = Sequential()
MobileNet_model.add(base_model)
MobileNet_model.add(Flatten())
MobileNet_model.add(Dense(100,activation=('softmax')))

```

```

early_stopping = EarlyStopping(min_delta = 0.001,patience =
                                20,restore_best_weights = True,verbose = 0)

```

```

# Compile

```

```

MobileNet_model.compile(optimizer = "adam" , loss =
                        'sparse_categorical_crossentropy' , metrics = ['accuracy'])

```



```

# Train
Mobile = MobileNet_model.fit(X_train_re, y_train, batch_size = 400,
                             epochs = 50, callbacks = [early_stopping], validation_data =
                             (X_test_re, y_test))

MobileNet_model.summary()

MobileNet_model.save('/content/drive/MyDrive/AML mini
project/mobile.h5')

# early_stopping = EarlyStopping(min_delta = 0.001, patience =
                             20, restore_best_weights = True, verbose = 2)
# model = MobileNet(input_shape=(32, 32, 3), alpha=1., weights=None,
                             classes=100)
# model.compile(optimizer=Adam(lr=0.002),
                             loss='categorical_crossentropy',
#                             metrics=['accuracy'])

# Mobile = model.fit(X_train_re, y_train, batch_size = 400, epochs =
                             50, callbacks = [early_stopping], validation_data =
                             (X_test_re, y_test))
# print(model.summary())

```

LeNet

reference: <https://towardsdatascience.com/going-beyond-99-mnist-handwritten-digits-recognition-cfff96337392>

```

train_path =
    "/content/drive/MyDrive/AMLminiproject/train100c5k_v2.pkl"
test_path =
    "/content/drive/MyDrive/AMLminiproject/test100c5k_nolabel.pkl"
# train_data = open(train_path, 'rb')
# test_data = open(test_path, 'rb')

# df_train = pd.read_pickle(train_path)
# df_test = pd.read_pickle(test_path)
df_train = pickle.load(open(train_path, 'rb'))
df_test = pickle.load(open(test_path, 'rb'))
# train_data = df['data'].values

# sampling
# sample = df_train.groupby('target', group_keys=False).apply(lambda x:
#                             x.sample(frac=0.3))
sample = df_train.copy()
sample['data'] = sample['data']/255

```

```

# X = np.array(list(sample['data'].values))
X = np.array([img[:, :, newaxis] for img in
              np.array(sample['data'].values)])
y = np.array(sample['target'].values)
df_train = None
df_test = None
sample = None
X.shape

(500000, 28, 28, 1)

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    random_state=42)
X = None
y = None
X_train.shape

(375000, 28, 28, 1)

X_train.shape

(375000, 28, 28, 1)

X_train_re = np.pad(X_train, ((0,0),(2,2),(2,2),(0,0)), 'constant')
X_test_re = np.pad(X_test, ((0,0),(2,2),(2,2),(0,0)), 'constant')
X_train = None
X_test = None
X_train_re.shape

(375000, 32, 32, 1)

y_train = to_categorical(y_train, 100)
y_test = to_categorical(y_test, 100)

```

LeNet architecture with additional hidden layers and batch normalization followed by them

```

model = Sequential([
    # Layer 1
    Conv2D(filters = 32, kernel_size = 5, strides = 1, activation =
            'relu', input_shape = (32,32,1),
            kernel_regularizer=regularizers.l1_l2(l1=0, l2=0.0005)),
    # Layer 2
    Conv2D(filters = 32, kernel_size = 5, strides = 1, use_bias=False),
    # Layer 3

```

```

BatchNormalization(),
# - - - - - #
Activation('relu'),
MaxPooling2D(pool_size = 2, strides = 2),
Dropout(0.25),
# - - - - - #
# Layer 3
Conv2D(filters = 64, kernel_size = 3, strides = 1, activation =
      'relu', kernel_regularizer=regularizers.l1_l2(l1=0,
      l2=0.0005)),
# Layer 4
Conv2D(filters = 64, kernel_size = 3, strides = 1, use_bias=False),
# Layer 5
BatchNormalization(),
# - - - - - #
Activation('relu'),
MaxPooling2D(pool_size = 2, strides = 2),
Dropout(0.25),
Flatten(),
# - - - - - #
# Layer 6
Dense(units = 256, use_bias=False),
# Layer 7
BatchNormalization(),
# - - - - - #
Activation('relu'),
# - - - - - #
# Layer 8
Dense(units = 128, use_bias=False),
# Layer 9
BatchNormalization(),
# - - - - - #
Activation('relu'),
# - - - - - #
# Layer 10
Dense(units = 84, use_bias=False),
# Layer 11
BatchNormalization(),
# - - - - - #
Activation('relu'),
Dropout(0.25),

```

```
# - - - - - #
# Output
Dense(units = 100, activation = 'softmax')
])

early_stopping = EarlyStopping(min_delta = 0.001, patience =
    20, restore_best_weights = True, verbose = 0)
model.compile(optimizer='adam', loss='categorical_crossentropy',
    metrics=['accuracy'])

# model.fit(X_train_re, y_train, batch_size = 400, epochs =
    50, callbacks = [early_stopping], validation_data = (X_test_re,
    y_test))
history = model.fit(X_train_re, y_train, epochs = 50, batch_size =
    100, callbacks = [early_stopping], validation_data =
    (X_test_re, y_test))
```

Epoch 1/50

3750/3750 [=====] - 54s 11ms/step - loss:
2.2783 - accuracy: 0.4449 - val_loss: 1.6516 - val_accuracy: 0.5801

Epoch 2/50

3750/3750 [=====] - 42s 11ms/step - loss:
1.6669 - accuracy: 0.5811 - val_loss: 1.4732 - val_accuracy: 0.6205

Epoch 3/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.5325 - accuracy: 0.6130 - val_loss: 1.3022 - val_accuracy: 0.6662

Epoch 4/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.4575 - accuracy: 0.6315 - val_loss: 1.2603 - val_accuracy: 0.6777

Epoch 5/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.4107 - accuracy: 0.6435 - val_loss: 1.2330 - val_accuracy: 0.6849

Epoch 6/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.3774 - accuracy: 0.6511 - val_loss: 1.2859 - val_accuracy: 0.6710

Epoch 7/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.3474 - accuracy: 0.6589 - val_loss: 1.1728 - val_accuracy: 0.6997

Epoch 8/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.3262 - accuracy: 0.6645 - val_loss: 1.1501 - val_accuracy: 0.7063

Epoch 9/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.3089 - accuracy: 0.6690 - val_loss: 1.2101 - val_accuracy: 0.6926

Epoch 10/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.2971 - accuracy: 0.6720 - val_loss: 1.2375 - val_accuracy: 0.6868

Epoch 11/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.2832 - accuracy: 0.6750 - val_loss: 1.1592 - val_accuracy: 0.7046

Epoch 12/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.2681 - accuracy: 0.6789 - val_loss: 1.1577 - val_accuracy: 0.7052

Epoch 13/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.2565 - accuracy: 0.6821 - val_loss: 1.1342 - val_accuracy: 0.7117

Epoch 14/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.2473 - accuracy: 0.6844 - val_loss: 1.1447 - val_accuracy: 0.7090

Epoch 15/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.2405 - accuracy: 0.6861 - val_loss: 1.1253 - val_accuracy: 0.7134

Epoch 16/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.2326 - accuracy: 0.6877 - val_loss: 1.1250 - val_accuracy: 0.7139

Epoch 17/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.2237 - accuracy: 0.6898 - val_loss: 1.0995 - val_accuracy: 0.7190

Epoch 18/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.2178 - accuracy: 0.6915 - val_loss: 1.1277 - val_accuracy: 0.7145

Epoch 19/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.2120 - accuracy: 0.6925 - val_loss: 1.0959 - val_accuracy: 0.7211

Epoch 20/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.2063 - accuracy: 0.6940 - val_loss: 1.0966 - val_accuracy: 0.7214

Epoch 21/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.2006 - accuracy: 0.6961 - val_loss: 1.1016 - val_accuracy: 0.7194

Epoch 22/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1943 - accuracy: 0.6963 - val_loss: 1.1236 - val_accuracy: 0.7140

Epoch 23/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.1898 - accuracy: 0.6979 - val_loss: 1.1178 - val_accuracy: 0.7157
Epoch 24/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1842 - accuracy: 0.6994 - val_loss: 1.0987 - val_accuracy: 0.7211
Epoch 25/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1826 - accuracy: 0.6997 - val_loss: 1.0938 - val_accuracy: 0.7225
Epoch 26/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.1776 - accuracy: 0.7012 - val_loss: 1.0939 - val_accuracy: 0.7221
Epoch 27/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1742 - accuracy: 0.7020 - val_loss: 1.0851 - val_accuracy: 0.7240
Epoch 28/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1715 - accuracy: 0.7027 - val_loss: 1.0933 - val_accuracy: 0.7220
Epoch 29/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1690 - accuracy: 0.7031 - val_loss: 1.0715 - val_accuracy: 0.7274
Epoch 30/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1622 - accuracy: 0.7045 - val_loss: 1.0896 - val_accuracy: 0.7241
Epoch 31/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1633 - accuracy: 0.7047 - val_loss: 1.0927 - val_accuracy: 0.7245
Epoch 32/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.1583 - accuracy: 0.7061 - val_loss: 1.0938 - val_accuracy: 0.7230
Epoch 33/50

3750/3750 [=====] - 32s 8ms/step - loss:
1.1554 - accuracy: 0.7065 - val_loss: 1.0898 - val_accuracy: 0.7233
Epoch 34/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1521 - accuracy: 0.7068 - val_loss: 1.0800 - val_accuracy: 0.7269
Epoch 35/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1514 - accuracy: 0.7075 - val_loss: 1.0962 - val_accuracy: 0.7238
Epoch 36/50

3750/3750 [=====] - 34s 9ms/step - loss:

1.1505 - accuracy: 0.7074 - val_loss: 1.1167 - val_accuracy: 0.7165
Epoch 37/50
3750/3750 [=====] - 32s 8ms/step - loss:
1.1466 - accuracy: 0.7085 - val_loss: 1.0998 - val_accuracy: 0.7219
Epoch 38/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1423 - accuracy: 0.7099 - val_loss: 1.0994 - val_accuracy: 0.7205
Epoch 39/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1412 - accuracy: 0.7095 - val_loss: 1.0968 - val_accuracy: 0.7220
Epoch 40/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1361 - accuracy: 0.7107 - val_loss: 1.0621 - val_accuracy: 0.7308
Epoch 41/50
3750/3750 [=====] - 32s 8ms/step - loss:
1.1384 - accuracy: 0.7110 - val_loss: 1.0644 - val_accuracy: 0.7309
Epoch 42/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1306 - accuracy: 0.7119 - val_loss: 1.0785 - val_accuracy: 0.7276
Epoch 43/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.1335 - accuracy: 0.7110 - val_loss: 1.0780 - val_accuracy: 0.7275
Epoch 44/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1282 - accuracy: 0.7127 - val_loss: 1.1323 - val_accuracy: 0.7139
Epoch 45/50
3750/3750 [=====] - 35s 9ms/step - loss:
1.1290 - accuracy: 0.7122 - val_loss: 1.0585 - val_accuracy: 0.7326
Epoch 46/50
3750/3750 [=====] - 32s 8ms/step - loss:
1.1273 - accuracy: 0.7128 - val_loss: 1.0745 - val_accuracy: 0.7280
Epoch 47/50
3750/3750 [=====] - 32s 9ms/step - loss:
1.1273 - accuracy: 0.7129 - val_loss: 1.0550 - val_accuracy: 0.7332
Epoch 48/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.1233 - accuracy: 0.7137 - val_loss: 1.0631 - val_accuracy: 0.7302
Epoch 49/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.1224 - accuracy: 0.7137 - val_loss: 1.0843 - val_accuracy: 0.7255

Epoch 50/50

3750/3750 [=====] - 40s 11ms/step - loss: 1.1189 - accuracy: 0.7147 - val_loss: 1.0673 - val_accuracy: 0.7298

```
model.save('/content/drive/MyDrive/AML mini project/lenet.h5')
```

```
save_path = '/content/drive/MyDrive/AML mini project/lenet.h5'
```

```
model.save(save_path, save_format='tf')
```

```
save_path = '/content/drive/MyDrive/AMLminiproject/lenet.h5'
```

```
model = tf.keras.models.load_model(save_path)
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 32)	832
conv2d_1 (Conv2D)	(None, 24, 24, 32)	25600
batch_normalization (Batch Normalization)	(None, 24, 24, 32)	128
activation (Activation)	(None, 24, 24, 32)	0
max_pooling2d (MaxPooling2D)	(None, 12, 12, 32)	0
dropout (Dropout)	(None, 12, 12, 32)	0
conv2d_2 (Conv2D)	(None, 10, 10, 64)	18496
conv2d_3 (Conv2D)	(None, 8, 8, 64)	36864
batch_normalization_1 (Batch Normalization)	(None, 8, 8, 64)	256
activation_1 (Activation)	(None, 8, 8, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 64)	0

dropout_1 (Dropout)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 256)	262144
batch_normalization_2 (Batch Normalization)	(None, 256)	1024
activation_2 (Activation)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32768
batch_normalization_3 (Batch Normalization)	(None, 128)	512
activation_3 (Activation)	(None, 128)	0
dense_2 (Dense)	(None, 84)	10752
batch_normalization_4 (Batch Normalization)	(None, 84)	336
activation_4 (Activation)	(None, 84)	0
dropout_2 (Dropout)	(None, 84)	0
dense_3 (Dense)	(None, 100)	8500

=====

Total params: 398,212

Trainable params: 397,084

Non-trainable params: 1,128

training for more 50 epochs i.e total of 100 epochs

```
early_stopping = EarlyStopping(min_delta = 0.001, patience =
    20, restore_best_weights = True, verbose = 0)
```

```
history = model.fit(X_train_re, y_train, epochs = 50, batch_size =  
100, callbacks = [early_stopping], validation_data =  
(X_test_re, y_test))
```

Epoch 1/50

3750/3750 [=====] - 47s 10ms/step - loss:
1.1188 - accuracy: 0.7156 - val_loss: 1.0998 - val_accuracy: 0.7250

Epoch 2/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1188 - accuracy: 0.7154 - val_loss: 1.0536 - val_accuracy: 0.7333

Epoch 3/50

3750/3750 [=====] - 36s 10ms/step - loss:
1.1176 - accuracy: 0.7152 - val_loss: 1.0578 - val_accuracy: 0.7320

Epoch 4/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.1138 - accuracy: 0.7159 - val_loss: 1.0859 - val_accuracy: 0.7266

Epoch 5/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1124 - accuracy: 0.7159 - val_loss: 1.0568 - val_accuracy: 0.7334

Epoch 6/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1119 - accuracy: 0.7168 - val_loss: 1.0853 - val_accuracy: 0.7261

Epoch 7/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.1092 - accuracy: 0.7173 - val_loss: 1.0525 - val_accuracy: 0.7336

Epoch 8/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.1099 - accuracy: 0.7166 - val_loss: 1.0811 - val_accuracy: 0.7275

Epoch 9/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.1091 - accuracy: 0.7174 - val_loss: 1.0553 - val_accuracy: 0.7328

Epoch 10/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.1065 - accuracy: 0.7175 - val_loss: 1.0673 - val_accuracy: 0.7310

Epoch 11/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.1051 - accuracy: 0.7182 - val_loss: 1.0560 - val_accuracy: 0.7336

Epoch 12/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1051 - accuracy: 0.7188 - val_loss: 1.0638 - val_accuracy: 0.7315

Epoch 13/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1011 - accuracy: 0.7195 - val_loss: 1.0767 - val_accuracy: 0.7279

Epoch 14/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.1019 - accuracy: 0.7186 - val_loss: 1.0484 - val_accuracy: 0.7347

Epoch 15/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1014 - accuracy: 0.7193 - val_loss: 1.0647 - val_accuracy: 0.7318

Epoch 16/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.0990 - accuracy: 0.7199 - val_loss: 1.0630 - val_accuracy: 0.7326

Epoch 17/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.0970 - accuracy: 0.7200 - val_loss: 1.0759 - val_accuracy: 0.7294

Epoch 18/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0971 - accuracy: 0.7206 - val_loss: 1.0858 - val_accuracy: 0.7267

Epoch 19/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.0967 - accuracy: 0.7205 - val_loss: 1.0699 - val_accuracy: 0.7298

Epoch 20/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.0963 - accuracy: 0.7207 - val_loss: 1.0471 - val_accuracy: 0.7356

Epoch 21/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.0961 - accuracy: 0.7209 - val_loss: 1.0690 - val_accuracy: 0.7310

Epoch 22/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0943 - accuracy: 0.7215 - val_loss: 1.0556 - val_accuracy: 0.7338

Epoch 23/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.0939 - accuracy: 0.7212 - val_loss: 1.0651 - val_accuracy: 0.7323

Epoch 24/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.0921 - accuracy: 0.7212 - val_loss: 1.0574 - val_accuracy: 0.7337

Epoch 25/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.0924 - accuracy: 0.7207 - val_loss: 1.0836 - val_accuracy: 0.7282

Epoch 26/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.0929 - accuracy: 0.7214 - val_loss: 1.0697 - val_accuracy: 0.7306
Epoch 27/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0901 - accuracy: 0.7218 - val_loss: 1.0572 - val_accuracy: 0.7345
Epoch 28/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.0885 - accuracy: 0.7227 - val_loss: 1.0650 - val_accuracy: 0.7323
Epoch 29/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0878 - accuracy: 0.7228 - val_loss: 1.0554 - val_accuracy: 0.7341
Epoch 30/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.0869 - accuracy: 0.7233 - val_loss: 1.0543 - val_accuracy: 0.7348
Epoch 31/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0873 - accuracy: 0.7232 - val_loss: 1.0550 - val_accuracy: 0.7342
Epoch 32/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.0871 - accuracy: 0.7227 - val_loss: 1.0521 - val_accuracy: 0.7358
Epoch 33/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0861 - accuracy: 0.7234 - val_loss: 1.0510 - val_accuracy: 0.7351
Epoch 34/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.0860 - accuracy: 0.7228 - val_loss: 1.0557 - val_accuracy: 0.7346
Epoch 35/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.0829 - accuracy: 0.7234 - val_loss: 1.0761 - val_accuracy: 0.7296
Epoch 36/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0814 - accuracy: 0.7245 - val_loss: 1.0670 - val_accuracy: 0.7321
Epoch 37/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0833 - accuracy: 0.7238 - val_loss: 1.0397 - val_accuracy: 0.7382
Epoch 38/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.0814 - accuracy: 0.7232 - val_loss: 1.0523 - val_accuracy: 0.7361
Epoch 39/50

3750/3750 [=====] - 34s 9ms/step - loss:

1.0819 - accuracy: 0.7242 - val_loss: 1.0573 - val_accuracy: 0.7328

Epoch 40/50

3750/3750 [=====] - 35s 9ms/step - loss:

1.0786 - accuracy: 0.7248 - val_loss: 1.0673 - val_accuracy: 0.7314

Epoch 41/50

3750/3750 [=====] - 32s 9ms/step - loss:

1.0809 - accuracy: 0.7241 - val_loss: 1.0668 - val_accuracy: 0.7325

Epoch 42/50

3750/3750 [=====] - 32s 9ms/step - loss:

1.0780 - accuracy: 0.7249 - val_loss: 1.0547 - val_accuracy: 0.7353

Epoch 43/50

3750/3750 [=====] - 34s 9ms/step - loss:

1.0772 - accuracy: 0.7252 - val_loss: 1.0579 - val_accuracy: 0.7352

Epoch 44/50

3750/3750 [=====] - 32s 9ms/step - loss:

1.0773 - accuracy: 0.7251 - val_loss: 1.0533 - val_accuracy: 0.7353

Epoch 45/50

3750/3750 [=====] - 33s 9ms/step - loss:

1.0762 - accuracy: 0.7252 - val_loss: 1.0487 - val_accuracy: 0.7371

Epoch 46/50

3750/3750 [=====] - 34s 9ms/step - loss:

1.0765 - accuracy: 0.7251 - val_loss: 1.0524 - val_accuracy: 0.7354

Epoch 47/50

3750/3750 [=====] - 33s 9ms/step - loss:

1.0748 - accuracy: 0.7252 - val_loss: 1.0475 - val_accuracy: 0.7381

Epoch 48/50

3750/3750 [=====] - 32s 9ms/step - loss:

1.0743 - accuracy: 0.7253 - val_loss: 1.0653 - val_accuracy: 0.7329

Epoch 49/50

3750/3750 [=====] - 34s 9ms/step - loss:

1.0733 - accuracy: 0.7259 - val_loss: 1.0532 - val_accuracy: 0.7353

Epoch 50/50

3750/3750 [=====] - 35s 9ms/step - loss:

1.0739 - accuracy: 0.7252 - val_loss: 1.0760 - val_accuracy: 0.7325

save_path = '/content/drive/MyDrive/AML mini project/lenet_100e.h5'

model.save(save_path, save_format='tf')

save_path = '/content/drive/MyDrive/AML mini project/lenet.h5'

model = tf.keras.models.load_model(save_path)

```
# we can see that the best fit is around 75 epochs, hence we will
    train the 50 epoch model for mre 25 epochs total of 75
history = model.fit(X_train_re, y_train, epochs = 25, batch_size =
    100, validation_data = (X_test_re, y_test))
```

```
-----
-----
NameError                                Traceback (most recent call
last)
```

```
<ipython-input-17-2bf365c77103> in <module>
```

```
1 # we can see that the best fit is around 75 epochs, hence we
will train the 50 epoch model for mre 25 epochs total of 75
```

```
----> 2 history = model.fit(X_train_re, y_train, epochs = 25,
batch_size = 100, validation_data = (X_test_re, y_test))
```

```
NameError: name 'model' is not defined
```

```
save_path = '/content/drive/MyDrive/AML mini project/lenet_75e.h5'
model.save(save_path, save_format='tf')
```

```
# df_test['data'] = df_test['data']/255
df_test.head()
```

	data
0	[[0, 0, 0, 4, 4, 0, 0, 3, 5, 2, 0, 0, 11, 0, 0...
1	[[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...
2	[[0, 3, 0, 0, 101, 128, 21, 0, 10, 4, 3, 2, 2,...
3	[[0, 12, 55, 101, 116, 98, 79, 66, 52, 39, 30,...
4	[[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,...

```
df_test['data'] = df_test['data']/255
```

```
df_test.head()
```

```
-----
-----
AttributeError                            Traceback (most recent call
last)
```

```
<ipython-input-27-711c74f5b217> in <module>
```

```
----> 1 df_test.head()
```

```

AttributeError: 'NoneType' object has no attribute 'head'

test_data = np.array([img[:, :, newaxis] for img in
                      np.array(df_test['data'].values)])
test_data.shape

(100000, 28, 28, 1)

X_test_data = np.pad(test_data, ((0,0),(2,2),(2,2),(0,0)), 'constant')
X_test_data.shape

(100000, 32, 32, 1)

save_path = '/content/drive/MyDrive/AML mini project/lenet_75e.h5'
model = tf.keras.models.load_model(save_path)

X_test_data[0].shape

(32, 32, 1)

y_pred = model.predict(X_test_data)

3125/3125 [=====] - 7s 2ms/step

results = np.argmax(y_pred,axis = 1)
results

array([77, 50, 80, ..., 63,  0, 22])

```

Trying modifying LeNet architecture

Increase batch size

```

model = Sequential([
# Layer 1
Conv2D(filters = 32, kernel_size = 5, strides = 1, activation =
      'relu', input_shape = (32,32,1),
      kernel_regularizer=regularizers.l1_l2(l1=0, l2=0.0005)),
# Layer 2
Conv2D(filters = 32, kernel_size = 5, strides = 1, use_bias=False),
# Layer 3
BatchNormalization(),
# - - - - - #
Activation('relu'),
MaxPooling2D(pool_size = 2, strides = 2),

```

```

Dropout(0.25),
# - - - - - #
# Layer 3
Conv2D(filters = 64, kernel_size = 3, strides = 1, activation =
      'relu', kernel_regularizer=regularizers.l1_l2(l1=0,
      l2=0.0005)),
# Layer 4
Conv2D(filters = 64, kernel_size = 3, strides = 1, use_bias=False),
# Layer 5
BatchNormalization(),
# - - - - - #
Activation('relu'),
MaxPooling2D(pool_size = 2, strides = 2),
Dropout(0.25),
Flatten(),
# - - - - - #
# Layer 6
Dense(units = 320, use_bias=False),
# Layer 7
BatchNormalization(),
# - - - - - #
Activation('relu'),
# - - - - - #
# Layer 8
Dense(units = 160, use_bias=False),
# Layer 9
BatchNormalization(),
# - - - - - #
Activation('relu'),
# - - - - - #
# Layer 10
Dense(units = 128, use_bias=False),
# Layer 11
BatchNormalization(),
# - - - - - #
Activation('relu'),
Dropout(0.25),
# - - - - - #
# Output
Dense(units = 100, activation = 'softmax')
])

```



```

early_stopping = EarlyStopping(min_delta = 0.001, patience =
                                20, restore_best_weights = True, verbose = 0)
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])

# model.fit(X_train_re, y_train, batch_size = 400, epochs =
            50, callbacks = [early_stopping], validation_data = (X_test_re,
            y_test))
history = model.fit(X_train_re, y_train, epochs = 50, batch_size =
                    150, callbacks = [early_stopping], validation_data =
                    (X_test_re, y_test))

```

Epoch 1/50

2500/2500 [=====] - 41s 13ms/step - loss: 2.2116 - accuracy: 0.4587 - val_loss: 1.5957 - val_accuracy: 0.5915

Epoch 2/50

2500/2500 [=====] - 32s 13ms/step - loss: 1.5810 - accuracy: 0.5981 - val_loss: 1.3743 - val_accuracy: 0.6450

Epoch 3/50

2500/2500 [=====] - 30s 12ms/step - loss: 1.4423 - accuracy: 0.6323 - val_loss: 1.2739 - val_accuracy: 0.6710

Epoch 4/50

2500/2500 [=====] - 30s 12ms/step - loss: 1.3668 - accuracy: 0.6505 - val_loss: 1.1954 - val_accuracy: 0.6905

Epoch 5/50

2500/2500 [=====] - 31s 12ms/step - loss: 1.3173 - accuracy: 0.6633 - val_loss: 1.2172 - val_accuracy: 0.6856

Epoch 6/50

2500/2500 [=====] - 29s 12ms/step - loss: 1.2796 - accuracy: 0.6715 - val_loss: 1.1713 - val_accuracy: 0.6970

Epoch 7/50

2500/2500 [=====] - 31s 12ms/step - loss: 1.2501 - accuracy: 0.6787 - val_loss: 1.1428 - val_accuracy: 0.7053

Epoch 8/50

2500/2500 [=====] - 31s 12ms/step - loss: 1.2269 - accuracy: 0.6847 - val_loss: 1.1322 - val_accuracy: 0.7077

Epoch 9/50

2500/2500 [=====] - 29s 12ms/step - loss: 1.2061 - accuracy: 0.6898 - val_loss: 1.1290 - val_accuracy: 0.7090

Epoch 10/50

2500/2500 [=====] - 31s 12ms/step - loss: 1.1898 - accuracy: 0.6943 - val_loss: 1.0969 - val_accuracy: 0.7168

Epoch 11/50

2500/2500 [=====] - 31s 12ms/step - loss:
1.1772 - accuracy: 0.6967 - val_loss: 1.1215 - val_accuracy: 0.7130
Epoch 12/50
2500/2500 [=====] - 29s 12ms/step - loss:
1.1615 - accuracy: 0.7003 - val_loss: 1.0929 - val_accuracy: 0.7197
Epoch 13/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.1515 - accuracy: 0.7030 - val_loss: 1.0917 - val_accuracy: 0.7193
Epoch 14/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.1431 - accuracy: 0.7050 - val_loss: 1.0779 - val_accuracy: 0.7221
Epoch 15/50
2500/2500 [=====] - 31s 13ms/step - loss:
1.1311 - accuracy: 0.7085 - val_loss: 1.0904 - val_accuracy: 0.7206
Epoch 16/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.1220 - accuracy: 0.7108 - val_loss: 1.0715 - val_accuracy: 0.7239
Epoch 17/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.1153 - accuracy: 0.7110 - val_loss: 1.0677 - val_accuracy: 0.7280
Epoch 18/50
2500/2500 [=====] - 32s 13ms/step - loss:
1.1077 - accuracy: 0.7131 - val_loss: 1.0531 - val_accuracy: 0.7295
Epoch 19/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.1004 - accuracy: 0.7154 - val_loss: 1.0938 - val_accuracy: 0.7196
Epoch 20/50
2500/2500 [=====] - 29s 12ms/step - loss:
1.0953 - accuracy: 0.7169 - val_loss: 1.0504 - val_accuracy: 0.7322
Epoch 21/50
2500/2500 [=====] - 29s 12ms/step - loss:
1.0891 - accuracy: 0.7175 - val_loss: 1.0610 - val_accuracy: 0.7271
Epoch 22/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.0817 - accuracy: 0.7195 - val_loss: 1.0567 - val_accuracy: 0.7311
Epoch 23/50
2500/2500 [=====] - 29s 12ms/step - loss:
1.0790 - accuracy: 0.7194 - val_loss: 1.0614 - val_accuracy: 0.7284
Epoch 24/50
2500/2500 [=====] - 29s 12ms/step - loss:

1.0723 - accuracy: 0.7222 - val_loss: 1.0620 - val_accuracy: 0.7291
Epoch 25/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.0685 - accuracy: 0.7226 - val_loss: 1.0496 - val_accuracy: 0.7308
Epoch 26/50
2500/2500 [=====] - 29s 12ms/step - loss:
1.0643 - accuracy: 0.7239 - val_loss: 1.0525 - val_accuracy: 0.7322
Epoch 27/50
2500/2500 [=====] - 29s 12ms/step - loss:
1.0615 - accuracy: 0.7247 - val_loss: 1.0622 - val_accuracy: 0.7284
Epoch 28/50
2500/2500 [=====] - 32s 13ms/step - loss:
1.0564 - accuracy: 0.7258 - val_loss: 1.0464 - val_accuracy: 0.7344
Epoch 29/50
2500/2500 [=====] - 31s 13ms/step - loss:
1.0523 - accuracy: 0.7264 - val_loss: 1.0464 - val_accuracy: 0.7329
Epoch 30/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0504 - accuracy: 0.7274 - val_loss: 1.0430 - val_accuracy: 0.7353
Epoch 31/50
2500/2500 [=====] - 32s 13ms/step - loss:
1.0470 - accuracy: 0.7276 - val_loss: 1.0809 - val_accuracy: 0.7252
Epoch 32/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0410 - accuracy: 0.7291 - val_loss: 1.0540 - val_accuracy: 0.7315
Epoch 33/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.0401 - accuracy: 0.7293 - val_loss: 1.0729 - val_accuracy: 0.7264
Epoch 34/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0361 - accuracy: 0.7307 - val_loss: 1.0624 - val_accuracy: 0.7309
Epoch 35/50
2500/2500 [=====] - 32s 13ms/step - loss:
1.0335 - accuracy: 0.7311 - val_loss: 1.0749 - val_accuracy: 0.7274
Epoch 36/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.0308 - accuracy: 0.7319 - val_loss: 1.0504 - val_accuracy: 0.7331
Epoch 37/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0275 - accuracy: 0.7329 - val_loss: 1.0430 - val_accuracy: 0.7358

Epoch 38/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.0257 - accuracy: 0.7324 - val_loss: 1.0396 - val_accuracy: 0.7358

Epoch 39/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.0230 - accuracy: 0.7340 - val_loss: 1.0398 - val_accuracy: 0.7369

Epoch 40/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.0236 - accuracy: 0.7333 - val_loss: 1.0351 - val_accuracy: 0.7371

Epoch 41/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.0172 - accuracy: 0.7352 - val_loss: 1.0561 - val_accuracy: 0.7349

Epoch 42/50
2500/2500 [=====] - 31s 13ms/step - loss:
1.0153 - accuracy: 0.7355 - val_loss: 1.0379 - val_accuracy: 0.7378

Epoch 43/50
2500/2500 [=====] - 29s 12ms/step - loss:
1.0141 - accuracy: 0.7355 - val_loss: 1.0442 - val_accuracy: 0.7354

Epoch 44/50
2500/2500 [=====] - 31s 12ms/step - loss:
1.0114 - accuracy: 0.7354 - val_loss: 1.0383 - val_accuracy: 0.7360

Epoch 45/50
2500/2500 [=====] - 29s 12ms/step - loss:
1.0093 - accuracy: 0.7362 - val_loss: 1.0398 - val_accuracy: 0.7371

Epoch 46/50
2500/2500 [=====] - 31s 13ms/step - loss:
1.0104 - accuracy: 0.7360 - val_loss: 1.0411 - val_accuracy: 0.7377

Epoch 47/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0043 - accuracy: 0.7384 - val_loss: 1.0775 - val_accuracy: 0.7279

Epoch 48/50
2500/2500 [=====] - 31s 13ms/step - loss:
1.0028 - accuracy: 0.7380 - val_loss: 1.0375 - val_accuracy: 0.7366

Epoch 49/50
2500/2500 [=====] - 32s 13ms/step - loss:
1.0010 - accuracy: 0.7383 - val_loss: 1.0387 - val_accuracy: 0.7370

Epoch 50/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0015 - accuracy: 0.7386 - val_loss: 1.0338 - val_accuracy: 0.7378

Trying to change neurons in hidden layers

```

model = Sequential([
# Layer 1
Conv2D(filters = 32, kernel_size = 5, strides = 1, activation =
    'relu', input_shape = (32,32,1),
    kernel_regularizer=regularizers.l1_l2(l1=0, l2=0.0005)),
# Layer 2
Conv2D(filters = 32, kernel_size = 5, strides = 1, use_bias=False),
# Layer 3
BatchNormalization(),
# - - - - - #
Activation('relu'),
MaxPooling2D(pool_size = 2, strides = 2),
Dropout(0.25),
# - - - - - #
# Layer 3
Conv2D(filters = 64, kernel_size = 3, strides = 1, activation =
    'relu', kernel_regularizer=regularizers.l1_l2(l1=0,
    l2=0.0005)),
# Layer 4
Conv2D(filters = 64, kernel_size = 3, strides = 1, use_bias=False),
# Layer 5
BatchNormalization(),
# - - - - - #
Activation('relu'),
MaxPooling2D(pool_size = 2, strides = 2),
Dropout(0.25),
Flatten(),
# - - - - - #
# Layer 6
Dense(units = 256, use_bias=False),
# Layer 7
BatchNormalization(),
# - - - - - #
Activation('relu'),
# - - - - - #
# Layer 10
Dense(units = 128, use_bias=False),
# Layer 11
BatchNormalization(),
# - - - - - #

```

```

Activation('relu'),
Dropout(0.25),
# - - - - - #
# Output
Dense(units = 100, activation = 'softmax')
])

early_stopping = EarlyStopping(min_delta = 0.001, patience =
    20, restore_best_weights = True, verbose = 0)
model.compile(optimizer='adam', loss='categorical_crossentropy',
    metrics=['accuracy'])

# model.fit(X_train_re, y_train, batch_size = 400, epochs =
    50, callbacks = [early_stopping], validation_data = (X_test_re,
    y_test))
history = model.fit(X_train_re, y_train, epochs = 50, batch_size =
    150, callbacks = [early_stopping], validation_data =
    (X_test_re, y_test))

```

Epoch 1/50

2500/2500 [=====] - 31s 12ms/step - loss: 2.2103 - accuracy: 0.4587 - val_loss: 1.6014 - val_accuracy: 0.5920

Epoch 2/50

2500/2500 [=====] - 29s 12ms/step - loss: 1.5995 - accuracy: 0.5932 - val_loss: 1.3521 - val_accuracy: 0.6519

Epoch 3/50

2500/2500 [=====] - 30s 12ms/step - loss: 1.4686 - accuracy: 0.6246 - val_loss: 1.3037 - val_accuracy: 0.6650

Epoch 4/50

2500/2500 [=====] - 28s 11ms/step - loss: 1.3992 - accuracy: 0.6420 - val_loss: 1.3229 - val_accuracy: 0.6599

Epoch 5/50

2500/2500 [=====] - 28s 11ms/step - loss: 1.3533 - accuracy: 0.6536 - val_loss: 1.2024 - val_accuracy: 0.6896

Epoch 6/50

2500/2500 [=====] - 28s 11ms/step - loss: 1.3196 - accuracy: 0.6615 - val_loss: 1.1897 - val_accuracy: 0.6926

Epoch 7/50

2500/2500 [=====] - 28s 11ms/step - loss: 1.2899 - accuracy: 0.6685 - val_loss: 1.1597 - val_accuracy: 0.7022

Epoch 8/50

2500/2500 [=====] - 30s 12ms/step - loss: 1.2709 - accuracy: 0.6730 - val_loss: 1.1717 - val_accuracy: 0.6966

Epoch 9/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.2518 - accuracy: 0.6780 - val_loss: 1.1402 - val_accuracy: 0.7075

Epoch 10/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.2366 - accuracy: 0.6815 - val_loss: 1.1636 - val_accuracy: 0.7006

Epoch 11/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.2245 - accuracy: 0.6843 - val_loss: 1.1111 - val_accuracy: 0.7146

Epoch 12/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.2137 - accuracy: 0.6867 - val_loss: 1.1004 - val_accuracy: 0.7170

Epoch 13/50
2500/2500 [=====] - 29s 11ms/step - loss:
1.2040 - accuracy: 0.6897 - val_loss: 1.1071 - val_accuracy: 0.7157

Epoch 14/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.1941 - accuracy: 0.6921 - val_loss: 1.1185 - val_accuracy: 0.7122

Epoch 15/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.1865 - accuracy: 0.6937 - val_loss: 1.1260 - val_accuracy: 0.7110

Epoch 16/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.1797 - accuracy: 0.6949 - val_loss: 1.1019 - val_accuracy: 0.7187

Epoch 17/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.1733 - accuracy: 0.6971 - val_loss: 1.0888 - val_accuracy: 0.7215

Epoch 18/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.1659 - accuracy: 0.6986 - val_loss: 1.1507 - val_accuracy: 0.7043

Epoch 19/50
2500/2500 [=====] - 29s 11ms/step - loss:
1.1612 - accuracy: 0.7001 - val_loss: 1.0709 - val_accuracy: 0.7264

Epoch 20/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.1545 - accuracy: 0.7015 - val_loss: 1.0786 - val_accuracy: 0.7232

Epoch 21/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.1490 - accuracy: 0.7021 - val_loss: 1.0697 - val_accuracy: 0.7243

Epoch 22/50

2500/2500 [=====] - 28s 11ms/step - loss:
1.1475 - accuracy: 0.7032 - val_loss: 1.1085 - val_accuracy: 0.7174
Epoch 23/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.1388 - accuracy: 0.7055 - val_loss: 1.0930 - val_accuracy: 0.7198
Epoch 24/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.1383 - accuracy: 0.7050 - val_loss: 1.0703 - val_accuracy: 0.7261
Epoch 25/50
2500/2500 [=====] - 29s 11ms/step - loss:
1.1337 - accuracy: 0.7065 - val_loss: 1.0836 - val_accuracy: 0.7227
Epoch 26/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.1295 - accuracy: 0.7073 - val_loss: 1.0606 - val_accuracy: 0.7294
Epoch 27/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.1251 - accuracy: 0.7086 - val_loss: 1.0879 - val_accuracy: 0.7215
Epoch 28/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.1241 - accuracy: 0.7083 - val_loss: 1.0568 - val_accuracy: 0.7302
Epoch 29/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.1203 - accuracy: 0.7097 - val_loss: 1.0652 - val_accuracy: 0.7285
Epoch 30/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.1182 - accuracy: 0.7096 - val_loss: 1.1191 - val_accuracy: 0.7119
Epoch 31/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.1151 - accuracy: 0.7108 - val_loss: 1.0513 - val_accuracy: 0.7318
Epoch 32/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.1112 - accuracy: 0.7114 - val_loss: 1.0834 - val_accuracy: 0.7229
Epoch 33/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.1097 - accuracy: 0.7122 - val_loss: 1.0634 - val_accuracy: 0.7276
Epoch 34/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.1073 - accuracy: 0.7128 - val_loss: 1.0734 - val_accuracy: 0.7251
Epoch 35/50
2500/2500 [=====] - 30s 12ms/step - loss:

1.1040 - accuracy: 0.7125 - val_loss: 1.0718 - val_accuracy: 0.7261
Epoch 36/50
2500/2500 [=====] - 29s 11ms/step - loss:
1.1036 - accuracy: 0.7127 - val_loss: 1.0558 - val_accuracy: 0.7292
Epoch 37/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.1017 - accuracy: 0.7144 - val_loss: 1.1079 - val_accuracy: 0.7157
Epoch 38/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0992 - accuracy: 0.7150 - val_loss: 1.0574 - val_accuracy: 0.7302
Epoch 39/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0958 - accuracy: 0.7158 - val_loss: 1.0509 - val_accuracy: 0.7309
Epoch 40/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0946 - accuracy: 0.7155 - val_loss: 1.1161 - val_accuracy: 0.7151
Epoch 41/50
2500/2500 [=====] - 29s 12ms/step - loss:
1.0942 - accuracy: 0.7159 - val_loss: 1.0506 - val_accuracy: 0.7329
Epoch 42/50
2500/2500 [=====] - 28s 11ms/step - loss:
1.0900 - accuracy: 0.7163 - val_loss: 1.0706 - val_accuracy: 0.7265
Epoch 43/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0888 - accuracy: 0.7164 - val_loss: 1.0460 - val_accuracy: 0.7339
Epoch 44/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0866 - accuracy: 0.7170 - val_loss: 1.0615 - val_accuracy: 0.7298
Epoch 45/50
2500/2500 [=====] - 29s 11ms/step - loss:
1.0852 - accuracy: 0.7171 - val_loss: 1.0690 - val_accuracy: 0.7258
Epoch 46/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0855 - accuracy: 0.7179 - val_loss: 1.0500 - val_accuracy: 0.7329
Epoch 47/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0812 - accuracy: 0.7188 - val_loss: 1.0631 - val_accuracy: 0.7292
Epoch 48/50
2500/2500 [=====] - 30s 12ms/step - loss:
1.0832 - accuracy: 0.7182 - val_loss: 1.0621 - val_accuracy: 0.7305

Epoch 49/50

2500/2500 [=====] - 28s 11ms/step - loss: 1.0792 - accuracy: 0.7190 - val_loss: 1.0559 - val_accuracy: 0.7306

Epoch 50/50

2500/2500 [=====] - 28s 11ms/step - loss: 1.0749 - accuracy: 0.7199 - val_loss: 1.0655 - val_accuracy: 0.7294

```
save_path = '/content/drive/MyDrive/AML mini
             project/lenet_50e_150b_2.h5'
```

```
model.save(save_path, save_format='tf')
```

Reduce batch size back to 100, with new architecture(modified hidden layers)

```
model = Sequential([
# Layer 1
Conv2D(filters = 32, kernel_size = 5, strides = 1, activation =
      'relu', input_shape = (32,32,1),
      kernel_regularizer=regularizers.l1_l2(l1=0, l2=0.0005)),
# Layer 2
Conv2D(filters = 32, kernel_size = 5, strides = 1, use_bias=False),
# Layer 3
BatchNormalization(),
# - - - - - #
Activation('relu'),
MaxPooling2D(pool_size = 2, strides = 2),
Dropout(0.25),
# - - - - - #
# Layer 3
Conv2D(filters = 64, kernel_size = 3, strides = 1, activation =
      'relu', kernel_regularizer=regularizers.l1_l2(l1=0,
      l2=0.0005)),
# Layer 4
Conv2D(filters = 64, kernel_size = 3, strides = 1, use_bias=False),
# Layer 5
BatchNormalization(),
# - - - - - #
Activation('relu'),
MaxPooling2D(pool_size = 2, strides = 2),
Dropout(0.25),
Flatten(),
# - - - - - #
# Layer 6
```

```

Dense(units = 256, use_bias=False),
# Layer 7
BatchNormalization(),
# - - - - - #
Activation('relu'),
# - - - - - #
# Layer 10
Dense(units = 128, use_bias=False),
# Layer 11
BatchNormalization(),
# - - - - - #
Activation('relu'),
Dropout(0.25),
# - - - - - #
# Output
Dense(units = 100, activation = 'softmax')
])

early_stopping = EarlyStopping(min_delta = 0.001, patience =
                                20, restore_best_weights = True, verbose = 0)
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])

# model.fit(X_train_re, y_train, batch_size = 400, epochs =
            50, callbacks = [early_stopping], validation_data = (X_test_re,
            y_test))
history = model.fit(X_train_re, y_train, epochs = 50, batch_size =
                    100, callbacks = [early_stopping], validation_data =
                    (X_test_re, y_test))

Epoch 1/50
3750/3750 [=====] - 43s 9ms/step - loss:
2.1532 - accuracy: 0.4711 - val_loss: 1.5326 - val_accuracy: 0.6105
Epoch 2/50
3750/3750 [=====] - 32s 8ms/step - loss:
1.6008 - accuracy: 0.5934 - val_loss: 1.3645 - val_accuracy: 0.6505
Epoch 3/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.4787 - accuracy: 0.6228 - val_loss: 1.2879 - val_accuracy: 0.6712
Epoch 4/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.4118 - accuracy: 0.6392 - val_loss: 1.2095 - val_accuracy: 0.6888
Epoch 5/50

```

3750/3750 [=====] - 33s 9ms/step - loss:
1.3690 - accuracy: 0.6508 - val_loss: 1.1964 - val_accuracy: 0.6921
Epoch 6/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.3377 - accuracy: 0.6578 - val_loss: 1.1704 - val_accuracy: 0.6999
Epoch 7/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.3142 - accuracy: 0.6641 - val_loss: 1.1578 - val_accuracy: 0.7024
Epoch 8/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.2944 - accuracy: 0.6688 - val_loss: 1.1562 - val_accuracy: 0.7045
Epoch 9/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.2781 - accuracy: 0.6729 - val_loss: 1.1708 - val_accuracy: 0.6998
Epoch 10/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.2610 - accuracy: 0.6775 - val_loss: 1.1177 - val_accuracy: 0.7135
Epoch 11/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.2495 - accuracy: 0.6789 - val_loss: 1.1286 - val_accuracy: 0.7114
Epoch 12/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.2407 - accuracy: 0.6816 - val_loss: 1.1106 - val_accuracy: 0.7158
Epoch 13/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.2291 - accuracy: 0.6844 - val_loss: 1.1114 - val_accuracy: 0.7157
Epoch 14/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.2231 - accuracy: 0.6861 - val_loss: 1.1042 - val_accuracy: 0.7178
Epoch 15/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.2136 - accuracy: 0.6883 - val_loss: 1.0939 - val_accuracy: 0.7197
Epoch 16/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.2075 - accuracy: 0.6899 - val_loss: 1.1110 - val_accuracy: 0.7162
Epoch 17/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1975 - accuracy: 0.6922 - val_loss: 1.1220 - val_accuracy: 0.7115
Epoch 18/50
3750/3750 [=====] - 32s 8ms/step - loss:

1.1949 - accuracy: 0.6927 - val_loss: 1.0900 - val_accuracy: 0.7219
Epoch 19/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1884 - accuracy: 0.6938 - val_loss: 1.0867 - val_accuracy: 0.7224
Epoch 20/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1837 - accuracy: 0.6952 - val_loss: 1.0910 - val_accuracy: 0.7215
Epoch 21/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1786 - accuracy: 0.6965 - val_loss: 1.0993 - val_accuracy: 0.7205
Epoch 22/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.1745 - accuracy: 0.6973 - val_loss: 1.0767 - val_accuracy: 0.7240
Epoch 23/50
3750/3750 [=====] - 32s 8ms/step - loss:
1.1700 - accuracy: 0.6984 - val_loss: 1.0681 - val_accuracy: 0.7272
Epoch 24/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1681 - accuracy: 0.6988 - val_loss: 1.1138 - val_accuracy: 0.7176
Epoch 25/50
3750/3750 [=====] - 32s 8ms/step - loss:
1.1628 - accuracy: 0.7006 - val_loss: 1.0797 - val_accuracy: 0.7258
Epoch 26/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1576 - accuracy: 0.7022 - val_loss: 1.0691 - val_accuracy: 0.7264
Epoch 27/50
3750/3750 [=====] - 32s 8ms/step - loss:
1.1582 - accuracy: 0.7014 - val_loss: 1.0788 - val_accuracy: 0.7257
Epoch 28/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.1535 - accuracy: 0.7026 - val_loss: 1.0734 - val_accuracy: 0.7254
Epoch 29/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1483 - accuracy: 0.7038 - val_loss: 1.0970 - val_accuracy: 0.7198
Epoch 30/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.1466 - accuracy: 0.7040 - val_loss: 1.0746 - val_accuracy: 0.7260
Epoch 31/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.1456 - accuracy: 0.7052 - val_loss: 1.0736 - val_accuracy: 0.7265

Epoch 32/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1435 - accuracy: 0.7048 - val_loss: 1.0736 - val_accuracy: 0.7256

Epoch 33/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1372 - accuracy: 0.7068 - val_loss: 1.0707 - val_accuracy: 0.7278

Epoch 34/50

3750/3750 [=====] - 31s 8ms/step - loss:
1.1367 - accuracy: 0.7071 - val_loss: 1.0758 - val_accuracy: 0.7268

Epoch 35/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1344 - accuracy: 0.7072 - val_loss: 1.0673 - val_accuracy: 0.7304

Epoch 36/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1315 - accuracy: 0.7078 - val_loss: 1.0574 - val_accuracy: 0.7305

Epoch 37/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1320 - accuracy: 0.7083 - val_loss: 1.0718 - val_accuracy: 0.7268

Epoch 38/50

3750/3750 [=====] - 32s 8ms/step - loss:
1.1276 - accuracy: 0.7085 - val_loss: 1.0619 - val_accuracy: 0.7291

Epoch 39/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1264 - accuracy: 0.7093 - val_loss: 1.0466 - val_accuracy: 0.7337

Epoch 40/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1259 - accuracy: 0.7098 - val_loss: 1.0781 - val_accuracy: 0.7258

Epoch 41/50

3750/3750 [=====] - 31s 8ms/step - loss:
1.1213 - accuracy: 0.7106 - val_loss: 1.0522 - val_accuracy: 0.7309

Epoch 42/50

3750/3750 [=====] - 31s 8ms/step - loss:
1.1213 - accuracy: 0.7106 - val_loss: 1.0706 - val_accuracy: 0.7276

Epoch 43/50

3750/3750 [=====] - 32s 8ms/step - loss:
1.1195 - accuracy: 0.7111 - val_loss: 1.0562 - val_accuracy: 0.7310

Epoch 44/50

3750/3750 [=====] - 31s 8ms/step - loss:
1.1164 - accuracy: 0.7110 - val_loss: 1.0514 - val_accuracy: 0.7319

Epoch 45/50

```

3750/3750 [=====] - 33s 9ms/step - loss:
1.1146 - accuracy: 0.7119 - val_loss: 1.0632 - val_accuracy: 0.7295
Epoch 46/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.1143 - accuracy: 0.7117 - val_loss: 1.0546 - val_accuracy: 0.7326
Epoch 47/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1110 - accuracy: 0.7132 - val_loss: 1.0881 - val_accuracy: 0.7243
Epoch 48/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1132 - accuracy: 0.7127 - val_loss: 1.0584 - val_accuracy: 0.7306
Epoch 49/50
3750/3750 [=====] - 32s 8ms/step - loss:
1.1096 - accuracy: 0.7133 - val_loss: 1.0570 - val_accuracy: 0.7328
Epoch 50/50
3750/3750 [=====] - 31s 8ms/step - loss:
1.1082 - accuracy: 0.7135 - val_loss: 1.0591 - val_accuracy: 0.7299

save_path = '/content/drive/MyDrive/AML mini
project/lenet_50e_100b_1.h5'
model.save(save_path, save_format='tf')

```

Adding one more hidden layer

```

model = Sequential([
# Layer 1
Conv2D(filters = 32, kernel_size = 5, strides = 1, activation =
'relu', input_shape = (32,32,1),
kernel_regularizer=regularizers.l1_l2(l1=0, l2=0.0005)),
# Layer 2
Conv2D(filters = 32, kernel_size = 5, strides = 1, use_bias=False),
# Layer 3
BatchNormalization(),
# - - - - - #
Activation('relu'),
MaxPooling2D(pool_size = 2, strides = 2),
Dropout(0.25),
# - - - - - #
# Layer 3
Conv2D(filters = 64, kernel_size = 3, strides = 1, activation =
'relu', kernel_regularizer=regularizers.l1_l2(l1=0,
l2=0.0005)),
# Layer 4

```

```

Conv2D(filters = 64, kernel_size = 3, strides = 1, use_bias=False),
# Layer 5
BatchNormalization(),
# - - - - - #
Activation('relu'),
MaxPooling2D(pool_size = 2, strides = 2),
Dropout(0.25),
Flatten(),
# - - - - - #
# Layer 6
Dense(units = 256, use_bias=False),
# Layer 7
BatchNormalization(),
# - - - - - #
Activation('relu'),
# - - - - - #
# Layer 8
Dense(units = 186, use_bias=False),
# Layer 9
BatchNormalization(),
# - - - - - #
Activation('relu'),
# - - - - - #
# Layer 10
Dense(units = 128, use_bias=False),
# Layer 11
BatchNormalization(),
# - - - - - #
Activation('relu'),
# - - - - - #
# Layer 12
Dense(units = 84, use_bias=False),
# Layer 13
BatchNormalization(),
# - - - - - #
Activation('relu'),
Dropout(0.25),
# - - - - - #
# Output

```



```

Dense(units = 100, activation = 'softmax')
])

early_stopping = EarlyStopping(min_delta = 0.001, patience =
                                20, restore_best_weights = True, verbose = 0)
model.compile(optimizer='adam', loss='categorical_crossentropy',
              metrics=['accuracy'])

# model.fit(X_train_re, y_train, batch_size = 400, epochs =
            50, callbacks = [early_stopping], validation_data = (X_test_re,
            y_test))
history = model.fit(X_train_re, y_train, epochs = 50, batch_size =
                    100, callbacks = [early_stopping], validation_data =
                    (X_test_re, y_test))

```

Epoch 1/50

3750/3750 [=====] - 38s 10ms/step - loss: 2.3009 - accuracy: 0.4401 - val_loss: 1.5902 - val_accuracy: 0.5925

Epoch 2/50

3750/3750 [=====] - 34s 9ms/step - loss: 1.6725 - accuracy: 0.5796 - val_loss: 1.4001 - val_accuracy: 0.6415

Epoch 3/50

3750/3750 [=====] - 34s 9ms/step - loss: 1.5302 - accuracy: 0.6150 - val_loss: 1.3544 - val_accuracy: 0.6550

Epoch 4/50

3750/3750 [=====] - 34s 9ms/step - loss: 1.4518 - accuracy: 0.6347 - val_loss: 1.3184 - val_accuracy: 0.6616

Epoch 5/50

3750/3750 [=====] - 34s 9ms/step - loss: 1.3979 - accuracy: 0.6478 - val_loss: 1.2264 - val_accuracy: 0.6860

Epoch 6/50

3750/3750 [=====] - 34s 9ms/step - loss: 1.3604 - accuracy: 0.6581 - val_loss: 1.1936 - val_accuracy: 0.6949

Epoch 7/50

3750/3750 [=====] - 34s 9ms/step - loss: 1.3326 - accuracy: 0.6646 - val_loss: 1.1751 - val_accuracy: 0.6995

Epoch 8/50

3750/3750 [=====] - 34s 9ms/step - loss: 1.3076 - accuracy: 0.6700 - val_loss: 1.1685 - val_accuracy: 0.7021

Epoch 9/50

3750/3750 [=====] - 33s 9ms/step - loss: 1.2868 - accuracy: 0.6753 - val_loss: 1.1796 - val_accuracy: 0.6988

Epoch 10/50

3750/3750 [=====] - 34s 9ms/step - loss:

1.2699 - accuracy: 0.6797 - val_loss: 1.2019 - val_accuracy: 0.6948
Epoch 11/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.2514 - accuracy: 0.6844 - val_loss: 1.1396 - val_accuracy: 0.7086
Epoch 12/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.2419 - accuracy: 0.6866 - val_loss: 1.1343 - val_accuracy: 0.7117
Epoch 13/50
3750/3750 [=====] - 35s 9ms/step - loss:
1.2315 - accuracy: 0.6892 - val_loss: 1.1484 - val_accuracy: 0.7083
Epoch 14/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.2181 - accuracy: 0.6925 - val_loss: 1.1194 - val_accuracy: 0.7162
Epoch 15/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.2098 - accuracy: 0.6940 - val_loss: 1.1274 - val_accuracy: 0.7134
Epoch 16/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.2026 - accuracy: 0.6949 - val_loss: 1.1141 - val_accuracy: 0.7166
Epoch 17/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.1927 - accuracy: 0.6979 - val_loss: 1.1511 - val_accuracy: 0.7071
Epoch 18/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.1854 - accuracy: 0.6999 - val_loss: 1.1378 - val_accuracy: 0.7115
Epoch 19/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1802 - accuracy: 0.7009 - val_loss: 1.1031 - val_accuracy: 0.7220
Epoch 20/50
3750/3750 [=====] - 33s 9ms/step - loss:
1.1708 - accuracy: 0.7042 - val_loss: 1.1022 - val_accuracy: 0.7199
Epoch 21/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.1653 - accuracy: 0.7049 - val_loss: 1.0740 - val_accuracy: 0.7275
Epoch 22/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.1634 - accuracy: 0.7052 - val_loss: 1.0928 - val_accuracy: 0.7229
Epoch 23/50
3750/3750 [=====] - 34s 9ms/step - loss:
1.1569 - accuracy: 0.7069 - val_loss: 1.0736 - val_accuracy: 0.7268

Epoch 24/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1520 - accuracy: 0.7077 - val_loss: 1.1143 - val_accuracy: 0.7171

Epoch 25/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1458 - accuracy: 0.7092 - val_loss: 1.0754 - val_accuracy: 0.7287

Epoch 26/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1428 - accuracy: 0.7105 - val_loss: 1.0879 - val_accuracy: 0.7255

Epoch 27/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1373 - accuracy: 0.7119 - val_loss: 1.0776 - val_accuracy: 0.7280

Epoch 28/50

3750/3750 [=====] - 32s 8ms/step - loss:
1.1334 - accuracy: 0.7128 - val_loss: 1.1104 - val_accuracy: 0.7176

Epoch 29/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1299 - accuracy: 0.7135 - val_loss: 1.0899 - val_accuracy: 0.7256

Epoch 30/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1269 - accuracy: 0.7135 - val_loss: 1.0791 - val_accuracy: 0.7266

Epoch 31/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1211 - accuracy: 0.7158 - val_loss: 1.0641 - val_accuracy: 0.7316

Epoch 32/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1197 - accuracy: 0.7160 - val_loss: 1.0664 - val_accuracy: 0.7300

Epoch 33/50

3750/3750 [=====] - 32s 8ms/step - loss:
1.1164 - accuracy: 0.7164 - val_loss: 1.0599 - val_accuracy: 0.7322

Epoch 34/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.1124 - accuracy: 0.7172 - val_loss: 1.0841 - val_accuracy: 0.7273

Epoch 35/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1100 - accuracy: 0.7181 - val_loss: 1.0623 - val_accuracy: 0.7311

Epoch 36/50

3750/3750 [=====] - 32s 9ms/step - loss:
1.1069 - accuracy: 0.7182 - val_loss: 1.0741 - val_accuracy: 0.7282

Epoch 37/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1046 - accuracy: 0.7196 - val_loss: 1.0906 - val_accuracy: 0.7242
Epoch 38/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.1035 - accuracy: 0.7191 - val_loss: 1.0509 - val_accuracy: 0.7352
Epoch 39/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0988 - accuracy: 0.7205 - val_loss: 1.0595 - val_accuracy: 0.7326
Epoch 40/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0981 - accuracy: 0.7206 - val_loss: 1.0623 - val_accuracy: 0.7338
Epoch 41/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.0954 - accuracy: 0.7216 - val_loss: 1.1022 - val_accuracy: 0.7211
Epoch 42/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0918 - accuracy: 0.7229 - val_loss: 1.0681 - val_accuracy: 0.7300
Epoch 43/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.0909 - accuracy: 0.7227 - val_loss: 1.0713 - val_accuracy: 0.7307
Epoch 44/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0886 - accuracy: 0.7242 - val_loss: 1.0763 - val_accuracy: 0.7291
Epoch 45/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.0854 - accuracy: 0.7228 - val_loss: 1.0581 - val_accuracy: 0.7340
Epoch 46/50

3750/3750 [=====] - 34s 9ms/step - loss:
1.0844 - accuracy: 0.7244 - val_loss: 1.0536 - val_accuracy: 0.7346
Epoch 47/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.0824 - accuracy: 0.7244 - val_loss: 1.0541 - val_accuracy: 0.7357
Epoch 48/50

3750/3750 [=====] - 33s 9ms/step - loss:
1.0825 - accuracy: 0.7242 - val_loss: 1.0521 - val_accuracy: 0.7361
Epoch 49/50

3750/3750 [=====] - 35s 9ms/step - loss:
1.0792 - accuracy: 0.7255 - val_loss: 1.0494 - val_accuracy: 0.7356
Epoch 50/50

```
3750/3750 [=====] - 35s 9ms/step - loss:
1.0761 - accuracy: 0.7261 - val_loss: 1.0434 - val_accuracy: 0.7368
```

```
save_path = '/content/drive/MyDrive/AML mini
project/lenet_50e_dense.h5'
model.save(save_path, save_format='tf')
```

```
X_train_re = None
```

Predictions

```
df_test = pickle.load(open(test_path, 'rb'))
df_test['data'] = df_test['data']/255
df_test.head()
```

	data
0	[[0.0, 0.0, 0.0, 0.01568627450980392, 0.015686...
1	[[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
2	[[0.0, 0.011764705882352941, 0.0, 0.0, 0.39607...
3	[[0.0, 0.047058823529411764, 0.215686274509803...
4	[[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...

```
test_data = np.array([img[:, :, newaxis] for img in
np.array(df_test['data'].values)])
```

```
test_data.shape
```

```
(100000, 28, 28, 1)
```

```
X_test_data = np.pad(test_data, ((0,0),(2,2),(2,2),(0,0)), 'constant')
```

```
X_test_data.shape
```

```
(100000, 32, 32, 1)
```

```
save_path = '/content/drive/MyDrive/AMLminiproject/lenet_75e.h5'
```

```
model = tf.keras.models.load_model(save_path)
```

```
y_pred = model.predict(X_test_data)
```

```
3125/3125 [=====] - 119s 38ms/step
```

```
results = np.argmax(y_pred, axis = 1)
```

```
results
```

```
array([77, 50, 80, ..., 63, 0, 22])

with open('/content/drive/MyDrive/AMLminiproject/project_rsheta.txt',
          'w') as file: # edit here as your username
    file.write('\n'.join(map(str, results)))
    file.flush()

y_pred = model.predict(X_test_re)

3907/3907 [=====] - 10s 2ms/step

results = np.argmax(y_pred,axis = 1)
results.shape

(125000,)

y_test_new = np.argmax(y_test,axis = 1)
y_test_new.shape

(125000,)

from sklearn.metrics import accuracy_score

accuracy_score(y_test_new, results)

0.731992

y_test_new[:5]

array([20, 39, 28, 26, 81])

results[:5]

array([51, 39, 28, 26, 81])
```