

Mini Project Report for P556 Applied Machine Learning

Name: Rushank Ghanshyam Sheta

IU ID: rsheta

Approach: Trying to solve it with reference to MNIST digit recognition/classification

Trials:

1. Transfer learning with MobileNet

Preprocessing steps:

1. Standardize input by dividing by 255
2. convert 28x28 matrix to 32x32(cubic resize)
3. Convert greyscale(1 channel) image to RGB(3 channel)
4. One hot encode labels

Performance: batchsize=400, epochs=50

Train acc: 95%

Test acc: 58%

Interpretation: The model is overfitting badly to training data and is unable to generalize well thats why test accuracy is very low compared to train accuracy

2. LeNet architecture(modified hidden layers with batch normalization)

Preprocessing steps:

1. Standardize input by dividing by 255
2. convert 28x28 matrix to 32x32(cubic resize)
3. Change dimensions of training array
4. One hot encode labels

Performance: batchsize=100, epochs=50, 75(best), 100

Train acc: 72%

Test acc: 73%

Interpretation: The model's test accuracy is higher than train accuracy which suggests that the overfitting problem has been solved.

Other experiments: Modifying the hidden layers(increasing/decreasing neurons), adding more hidden layers, changing the batch size does not result in a significant increase in test accuracy when compared to the LeNet Model with batchsize=100, epochs=75, moreover all of the experiments with LeNet results in comparable accuracies.

Conclusion: After several experiments, I found out that LeNet with the below given model summary outperforms every other model tested while trained for 75 epochs and batchsize of 100. In future we can experiment with batch normalization(I think this is what helped improving the score) and adding Regularization.

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 32)	832
conv2d_1 (Conv2D)	(None, 24, 24, 32)	25600
batch_normalization (Batch Normalization)	(None, 24, 24, 32)	128
activation (Activation)	(None, 24, 24, 32)	0
max_pooling2d (MaxPooling2D)	(None, 12, 12, 32)	0
dropout (Dropout)	(None, 12, 12, 32)	0
conv2d_2 (Conv2D)	(None, 10, 10, 64)	18496
conv2d_3 (Conv2D)	(None, 8, 8, 64)	36864
batch_normalization_1 (Batch Normalization)	(None, 8, 8, 64)	256
activation_1 (Activation)	(None, 8, 8, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 64)	0
dropout_1 (Dropout)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 256)	262144
batch_normalization_2 (Batch Normalization)	(None, 256)	1024
activation_2 (Activation)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32768
batch_normalization_3 (Batch Normalization)	(None, 128)	512
activation_3 (Activation)	(None, 128)	0
dense_2 (Dense)	(None, 84)	10752
batch_normalization_4 (Batch Normalization)	(None, 84)	336
activation_4 (Activation)	(None, 84)	0
dropout_2 (Dropout)	(None, 84)	0
dense_3 (Dense)	(None, 100)	8500
=====		
Total params: 398,212		
Trainable params: 397,084		
Non-trainable params: 1,128		

Epoch 75/100

loss: 1.0929 - accuracy: 0.7214 - val_loss: 1.0697 - val_accuracy: 0.7306

References:

- 1)<https://www.kaggle.com/code/sabarish2611/alexnet-vs-mobilenet-using-mnist-data#AlexNet-vs-MobileNet:-Comparing-the-performance>
- 2)<https://towardsdatascience.com/going-beyond-99-mnist-handwritten-digits-recognition-cfff96337392>